


AWS
re:Start
CHALLENGE LAB

Build and Access an RDS Server



WEEK 7





Overview

Amazon Relational Database Service (Amazon RDS) makes it easy to set up, operate, and scale a relational database in the cloud. It provides cost-efficient and resizable capacity while managing time-consuming database administration tasks, which allows you to focus on your applications and business.

Your Challenge

- Launch an Amazon RDS DB instance using either Amazon Aurora Provisioned DB or MySQL database engines.
- Connect (SSH) to the LinuxServer
- Install a MySQL client, and use it to connect to your db.
- Create a table RESTART with the following columns:
 - Student ID (Number),
 - Student Name,
 - Restart City,
 - Graduation Date (Date Time)
- Insert 10 sample rows into this table
- Select all rows from this table
- Create a table CLOUD_PRACTITIONER with the following columns:
 - Student ID (Number)
 - Certification Date (Date Time)
- Insert 5 sample rows into this table
- Select all rows from this table
- Perform an inner join between the 2 tables created above and display student ID, Student Name, Certification Date

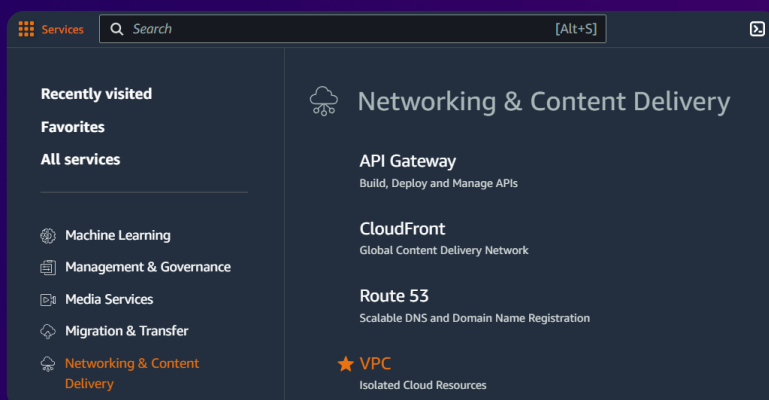


Task 1

Create a Security Group for the RDS DB Instance


Step 1: Access the VPC management console

Open the AWS Management Console, and select VPC.



Step 2: Create security group

Navigate to the **Security Groups** section, and select [Create security group](#).





Task 1

Create a Security Group for the RDS DB Instance

Step 3: Basic details

In the **Basic details** section, configure the DB Security Group using the following settings.

Basic details

Security group name [Info](#)

DB Security Group

Description [Info](#)

Permit access from Web Security Group

VPC [Info](#)


vpc-01cfc5333d09bd537 (Lab VPC)

Step 4: Inbound rules

In the **Inbound rules** section, configure the DB Security Group to permit inbound traffic on port 3306 from any EC2 instance that is associated with the Web Security Group.

Inbound rules [Info](#)

Type Info	Protocol Info	Port range Info	Source Info	Description - optional Info	
MySQL/Aurora	TCP	3306	Custom	sg-02a39c72741c7b2ef	<div>sg-02a39c72741c7b2ef X</div> <div>sg-02a39c72741c7b2ef X</div>
					Delete

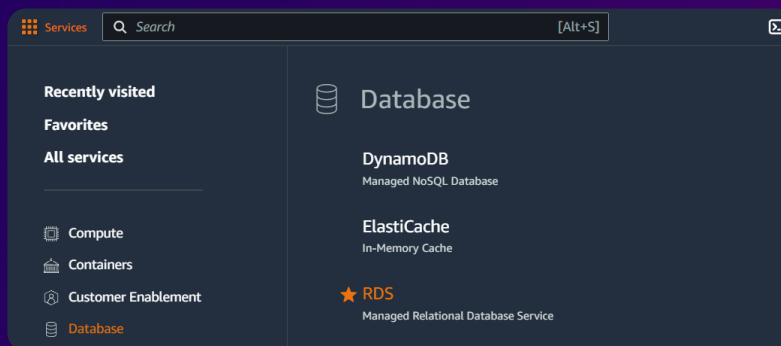


Task 2

Create a DB Subnet Group

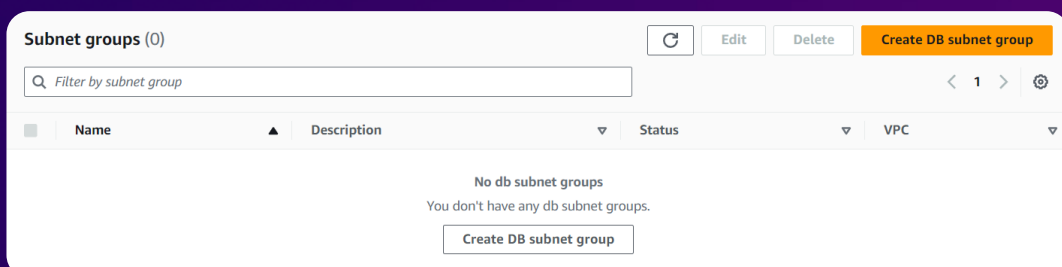
Step 1: Access the RDS database service

In the AWS Management Console, select RDS.



Step 2: Create DB subnet group

Navigate to the **Subnet groups** section, and select [Create DB subnet group](#).





Task 2

Create a DB Subnet Group

Step 3: Subnet group details

In the **Subnet group details** section, configure the DB Subnet Group using the following settings.

Subnet group details

Name

DB Subnet Group

Description

DB Subnet Group

VPC

Lab VPC (vpc-01cfc5333d09bd537)

Step 4: Add subnets

In the **Add subnets** section, configure the following settings.

Add subnets

Availability Zones

Choose the Availability Zones that include the subnets you want to add.

us-west-2a

us-west-2b

Subnets

Choose the subnets that you want to add. The list includes the subnets in the selected Availability Zones.

subnet-0581a875a79d996c8 (10.0.3.0/24)

subnet-0fb843e1120f98796 (10.0.1.0/24)

Subnets selected (2)

Availability zone	Subnet ID	CIDR block
us-west-2b	subnet-0581a875a79d996c8	10.0.3.0/24
us-west-2a	subnet-0fb843e1120f98796	10.0.1.0/24




Task 3

Launch an Amazon RDS DB instance

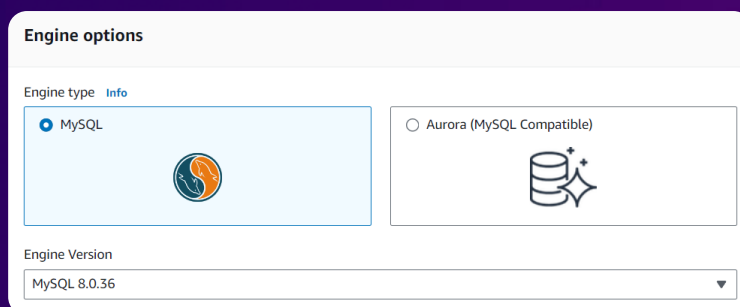
Step 1: Create database


Navigate to the Databases section, and select [Create database](#).



Step 2: Engine options

In the **Engine options** section, for Engine type, choose [MySQL](#), for Engine version, choose the latest version.





Task 3

Launch an Amazon RDS DB instance

Step 3: Templates

In the **Templates** section, choose [Free tier](#).

Templates
Choose a sample template to meet your use case.

☐ **Production**
Use defaults for high availability and fast, consistent performance.

☐ **Dev/Test**
This instance is intended for development use outside of a production environment.

☒ **Free tier**
Use RDS Free Tier to develop new applications, test existing applications, or gain hands-on experience with Amazon RDS.
[Info](#)

Step 4: Availability and durability

Notice that when you selected the [Free tier](#) template, the [Single DB instance](#) option was selected and locked as the default deployment option in the **Availability and durability** section.

Availability and durability

Deployment options [Info](#)
The deployment options below are limited to those supported by the engine you selected above.

☐ **Multi-AZ DB Cluster**
Creates a DB cluster with a primary DB instance and two readable standby DB instances, with each DB instance in a different Availability Zone (AZ). Provides high availability, data redundancy and increases capacity to serve read workloads.

☐ **Multi-AZ DB instance (not supported for Multi-AZ DB cluster snapshot)**
Creates a primary DB instance and a standby DB instance in a different AZ. Provides high availability and data redundancy, but the standby DB instance doesn't support connections for read workloads.

☒ **Single DB instance (not supported for Multi-AZ DB cluster snapshot)**
Creates a single DB instance with no standby DB instances.



Task 3

Launch an Amazon RDS DB instance

Step 5: Settings

In the **Settings** section, configure the following parameters.

Settings

DB instance identifier [Info](#)
Type a name for your DB instance.

▼ Credentials Settings

Master username [Info](#)
Type a login ID for the master user of your DB instance.

Master password [Info](#)

Confirm master password [Info](#)

Step 6: Instance configuration

In the **Instance configuration** section, for DB instance class, configure the following settings.

Instance configuration
The DB instance configuration options below are limited to those supported by the engine that you selected above.

DB instance class [Info](#)

☐ Standard classes (includes m classes)
☐ Memory optimized classes (includes r and x classes)
☒ Burstable classes (includes t classes)

2 vCPUs 1 GiB RAM Network: 2,085 Mbps



Task 3

Launch an Amazon RDS DB instance

Step 7: Storage

In the **Storage** section, for Storage type, select **General Purpose SSD (gp2)**.

The screenshot shows the 'Storage' configuration panel in the AWS console. It has a title bar 'Storage' with an 'Info' link. Below the title bar, it says 'Storage type Info' and 'Provisioned IOPS SSD (io2) storage volumes are now available.' There is a dropdown menu currently showing 'General Purpose SSD (gp2)' with a small downward arrow. Below the dropdown, it says 'Baseline performance determined by volume size'.

Step 8: Connectivity

In the **Connectivity** section, configure the following settings.

The screenshot shows the 'Connectivity' configuration panel in the AWS console. It has a title bar 'Connectivity Info' with a refresh icon. Below the title bar, it says 'Virtual private cloud (VPC) Info' and 'Choose the VPC. The VPC defines the virtual networking environment for this DB instance.' There is a dropdown menu showing 'Lab VPC (vpc-01cfc5333d09bd537)' with '4 Subnets, 2 Availability Zones' below it. Below that, it says 'DB subnet group Info' and 'Choose the DB subnet group. The DB subnet group defines which subnets and IP ranges the DB instance can use in the VPC that you selected.' There is a dropdown menu showing 'db subnet group' with '2 Subnets, 2 Availability Zones' below it. Below that, it says 'VPC security group (firewall) Info' and 'Choose one or more VPC security groups to allow access to your database. Make sure that the security group rules allow the appropriate incoming traffic.' There are two buttons: 'Choose existing' (selected) with 'Choose existing VPC security groups' below it, and 'Create new' with 'Create new VPC security group' below it. Below these buttons, it says 'Existing VPC security groups' and there is a dropdown menu showing 'Choose one or more options'. At the bottom, there is a button 'DB Security Group' with an 'X' icon.



Task 3

Launch an Amazon RDS DB instance

Step 9: Monitoring

In the **Monitoring** section, for Additional configuration, uncheck [Enable Enhanced Monitoring](#).

Monitoring

▼ Additional configuration

Enhanced Monitoring

☐ Enable Enhanced Monitoring

Enabling Enhanced Monitoring metrics are useful when you want to see how different processes or threads use the CPU.

Step 10: Review database creation

Verify the availability of the [challenge-lab-db](#) database and take note of its endpoint in the **Connectivity & Security** section.

Databases (1)

Group resources

Modify

Actions

Restore from S3

Create database

Filter by databases

< 1 > ⚙

<div><input type="checkbox"/></div> DB identifier	Status	Role	Engine	Region & AZ	Size	VPC
<div><div></div><div>challenge-lab-db</div></div>	<div><div></div><div>Available</div></div>	Instance	MySQL Community	us-west-2b	db.t3.micro	vpc-01cfc5333d09bd537



Use SSH to connect to the Linux Server


In the AWS Management Console, select the LinuxServer EC2 instance and make note of the **Public IPv4 address**.

Download the **private key file** [labsuser.pem](#). Change to the Downloads directory and modify the permissions on the key to be read-only (`r-----`).

Connect to the Linux Server using SSH

Establish a connection to the LinuxServer EC2 instance using the `ssh` command, the key and the instance's public IPv4 address.

```
support@HP-Pavilion-Laptop:~/Downloads$ ssh -i labsuser.pem ec2-user@52.40.199.25
The authenticity of host '52.40.199.25 (52.40.199.25)' can't be established.
ED25519 key fingerprint is SHA256:0BWIVAFxQPr2kYRdKvXdfjX8X0eJ4PFmF1g3mdB5Uvw.
This key is not known by any other names
Are you sure you want to continue connecting (yes/no/[fingerprint])? yes
Warning: Permanently added '52.40.199.25' (ED25519) to the list of known hosts.
```



```
#
#####      Amazon Linux 2
#####\
\###|      AL2 End of Life is 2025-06-30.
 \#/
  V-' '->

A newer version of Amazon Linux is available!

Amazon Linux 2023, GA and supported until 2028-03-15.
https://aws.amazon.com/linux/amazon-linux-2023/

[ec2-user@ip-10-0-2-249 ~]$
```



Task 5

Configure the EC2 instance to connect to the DB

Step 1: Install the DB client

Run the command `sudo yum install mariadb -y` to install the MariaDB client.

```
[ec2-user@ip-10-0-2-249 ~]$ sudo yum install mariadb -y
Loaded plugins: extras_suggestions, langpacks, priorities, update-motd
amzn2-core
Resolving Dependencies
--> Running transaction check
--> Package mariadb.x86_64 1:5.5.68-1.amzn2.0.1 will be installed
--> Finished Dependency Resolution
```

Step 2: Connect to the database

After installing a MySQL client, run the following command to connect to the database.

```
[ec2-user@ip-10-0-2-249 ~]$ mysql -u admin --password='lab-password' -h challenge-lab-db.chosw0e4egpw.us-west-2.rds.amazonaws.com
Welcome to the MariaDB monitor.  Commands end with ; or \g.
Your MySQL connection id is 27
Server version: 8.0.36 Source distribution

Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

MySQL [(none)]>
```



Task 6

Interact With Your DB

Step 1: Create a database

CREATE the **challenge_lab** database and switch to it.

```
MySQL [(none)]> CREATE DATABASE challenge_lab;
Query OK, 1 row affected (0.00 sec)

MySQL [(none)]> USE challenge_lab;
Database changed
MySQL [challenge_lab]>
```

Step 2: Create the **RESTART** table

CREATE a table **RESTART** with the following columns: Student ID (Number), Student Name, Restart City, Graduation Date (Date Time).

```
MySQL [challenge_lab]> CREATE TABLE RESTART (
->   `Student ID` INT PRIMARY KEY,
->   `Student Name` VARCHAR(255),
->   `Restart City` VARCHAR(255),
->   `Graduation Date` DATETIME
-> );
Query OK, 0 rows affected (0.02 sec)

MySQL [challenge_lab]>
```



Task 6

Interact With Your DB

Step 3: Insert sample rows

INSERT 10 sample rows into the RESTART table.

```
MySQL [challenge_lab]> INSERT INTO RESTART (`Student ID`, `Student Name`, `Restart City`, `Graduation Date`) VALUES
-> (1, 'John Doe', 'New York', '2023-05-15 10:00:00'),
-> (2, 'Jane Smith', 'Los Angeles', '2023-06-20 09:30:00'),
-> (3, 'Michael Johnson', 'Chicago', '2023-07-10 11:15:00'),
-> (4, 'Emily Davis', 'San Francisco', '2023-08-05 14:45:00'),
-> (5, 'David Brown', 'Houston', '2023-09-12 12:00:00'),
-> (6, 'Jessica Wilson', 'Miami', '2023-10-18 13:30:00'),
-> (7, 'Matthew Miller', 'Dallas', '2023-11-25 08:45:00'),
-> (8, 'Olivia Anderson', 'Seattle', '2023-12-10 09:00:00'),
-> (9, 'Ethan Martinez', 'Atlanta', '2024-01-05 10:30:00'),
-> (10, 'Ava Taylor', 'Denver', '2024-02-20 11:45:00');
Query OK, 10 rows affected (0.00 sec)
Records: 10  Duplicates: 0  Warnings: 0


MySQL [challenge_lab]>
```

Step 4: Select all rows

SELECT all rows from the RESTART table.

```
MySQL [challenge_lab]> SELECT * FROM RESTART;
+-----+-----+-----+-----+
| Student ID | Student Name | Restart City | Graduation Date |
+-----+-----+-----+-----+
| 1 | John Doe | New York | 2023-05-15 10:00:00 |
| 2 | Jane Smith | Los Angeles | 2023-06-20 09:30:00 |
| 3 | Michael Johnson | Chicago | 2023-07-10 11:15:00 |
| 4 | Emily Davis | San Francisco | 2023-08-05 14:45:00 |
| 5 | David Brown | Houston | 2023-09-12 12:00:00 |
| 6 | Jessica Wilson | Miami | 2023-10-18 13:30:00 |
| 7 | Matthew Miller | Dallas | 2023-11-25 08:45:00 |
| 8 | Olivia Anderson | Seattle | 2023-12-10 09:00:00 |
| 9 | Ethan Martinez | Atlanta | 2024-01-05 10:30:00 |
| 10 | Ava Taylor | Denver | 2024-02-20 11:45:00 |
+-----+-----+-----+-----+
10 rows in set (0.00 sec)

MySQL [challenge_lab]>
```



Task 6

Interact With Your DB

Step 5: Create the **CLOUD_PRACTITIONER** table

CREATE a table CLOUD_PRACTITIONER with the following columns: Student ID (Number), Certification Date (Date Time).

```
MySQL [challenge_lab]> CREATE TABLE CLOUD_PRACTITIONER (  
->   `Student ID` INT PRIMARY KEY,  
->   `Certification Date` DATETIME  
-> );  
Query OK, 0 rows affected (0.02 sec)  
  
MySQL [challenge_lab]>
```

Step 6: Insert sample rows

INSERT 5 sample rows into the CLOUD_PRACTITIONER table.

```
MySQL [challenge_lab]> INSERT INTO CLOUD_PRACTITIONER (`Student ID`, `Certification Date`) VALUES  
-> (1, '2023-05-10 10:00:00'),  
-> (2, '2023-06-15 09:30:00'),  
-> (3, '2023-07-05 11:15:00'),  
-> (4, '2023-07-31 14:45:00'),  
-> (5, '2023-09-07 12:00:00');  
Query OK, 5 rows affected (0.00 sec)  
Records: 5  Duplicates: 0  Warnings: 0  
  
MySQL [challenge_lab]>
```




Task 6

Interact With Your DB

Step 7: Select all rows

SELECT all rows from the CLOUD_PRACTITIONER table.

```
MySQL [challenge_lab]> SELECT * FROM CLOUD_PRACTITIONER;
+-----+-----+
| Student ID | Certification Date |
+-----+-----+
| 1 | 2023-05-10 10:00:00 |
| 2 | 2023-06-15 09:30:00 |
| 3 | 2023-07-05 11:15:00 |
| 4 | 2023-07-31 14:45:00 |
| 5 | 2023-09-07 12:00:00 |
+-----+-----+
5 rows in set (0.00 sec)

MySQL [challenge_lab]>
```

Step 8: Perform an inner join

Perform an **INNER JOIN** between the 2 tables created above and display Student ID, Student Name, Certification Date.

```
MySQL [challenge_lab]> SELECT RESTART.`Student ID`, `Student Name`, `Certification Date`
-> FROM RESTART
-> INNER JOIN CLOUD_PRACTITIONER
-> ON RESTART.`Student ID` = CLOUD_PRACTITIONER.`Student ID`;
+-----+-----+-----+
| Student ID | Student Name | Certification Date |
+-----+-----+-----+
| 1 | John Doe | 2023-05-10 10:00:00 |
| 2 | Jane Smith | 2023-06-15 09:30:00 |
| 3 | Michael Johnson | 2023-07-05 11:15:00 |
| 4 | Emily Davis | 2023-07-31 14:45:00 |
| 5 | David Brown | 2023-09-07 12:00:00 |
+-----+-----+-----+
5 rows in set (0.00 sec)

MySQL [challenge_lab]>
```



Conclusions

Amazon Relational Databases

Amazon Relational Databases offer scalable and reliable solutions for managing structured data, catering to diverse business needs.

Amazon RDS DB Instances

Amazon RDS DB Instances provide flexible configurations and high availability options, ensuring continuous access to databases.

Permitting connections to a DB instance

Permitting connections to a DB instance allows seamless communication between applications and databases, facilitating real-time data interactions.

DB Subnet Groups

DB Subnet Groups enable secure networking configurations, ensuring data privacy and compliance with regulatory requirements.

Interacting with a Database

Interacting with a database through applications or query tools enables data retrieval, updates, and analysis, empowering informed decision-making and efficient data management.



Cristhian Becerra



[cristhian-becerra-espinoza](#)



+51 951 634 354



cristhianbecerra99@gmail.com



Lima, Peru

